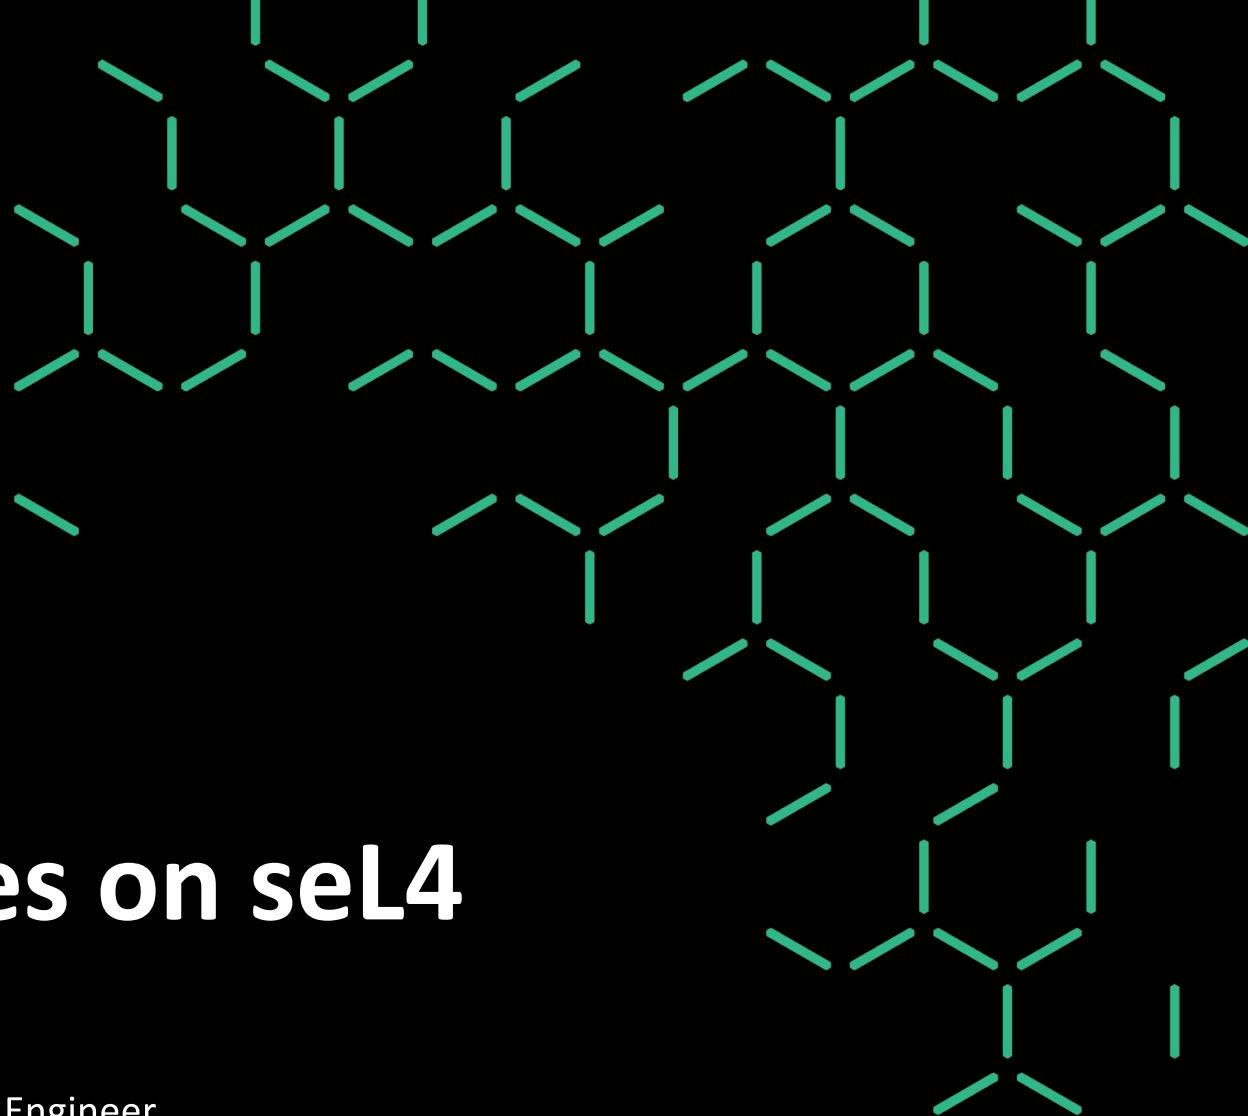




DATA
61



Using Devices on seL4

Stephen Sherratt

| Research Engineer

August 2016

www.data61.csiro.au

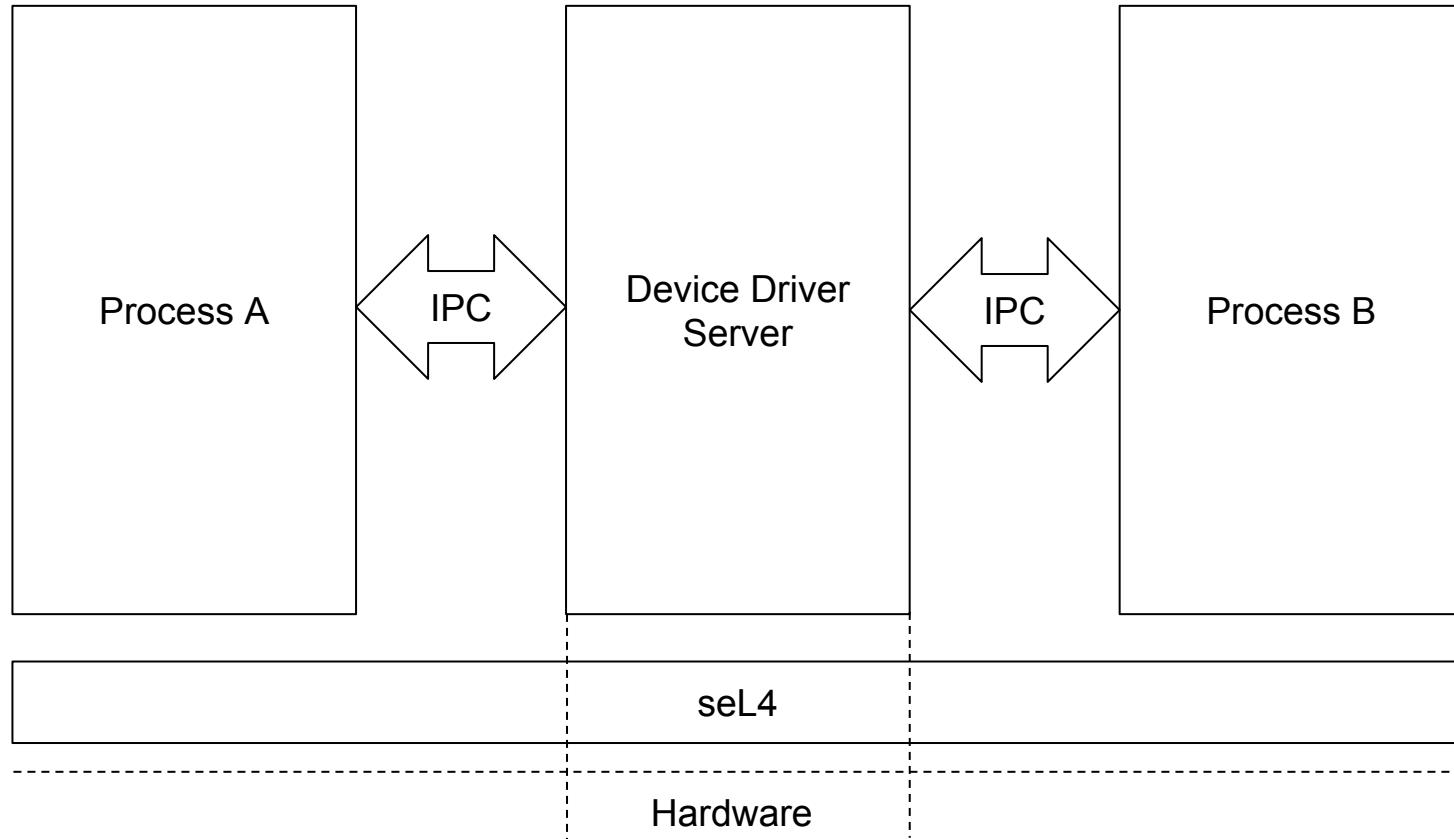


Overview

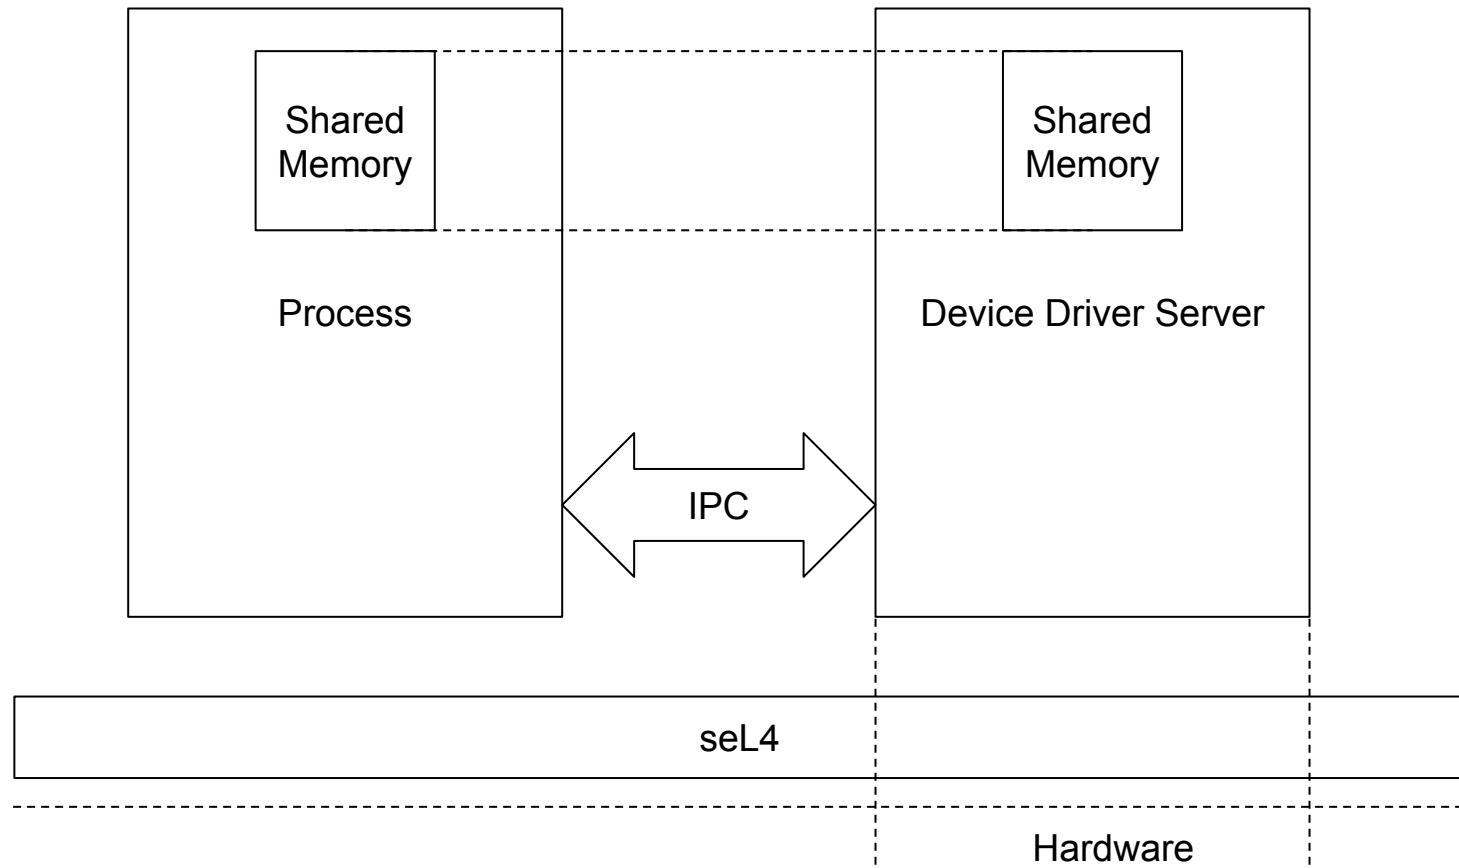


- Driver Model
- Interacting with Hardware
- Using Drivers
- Writing Drivers

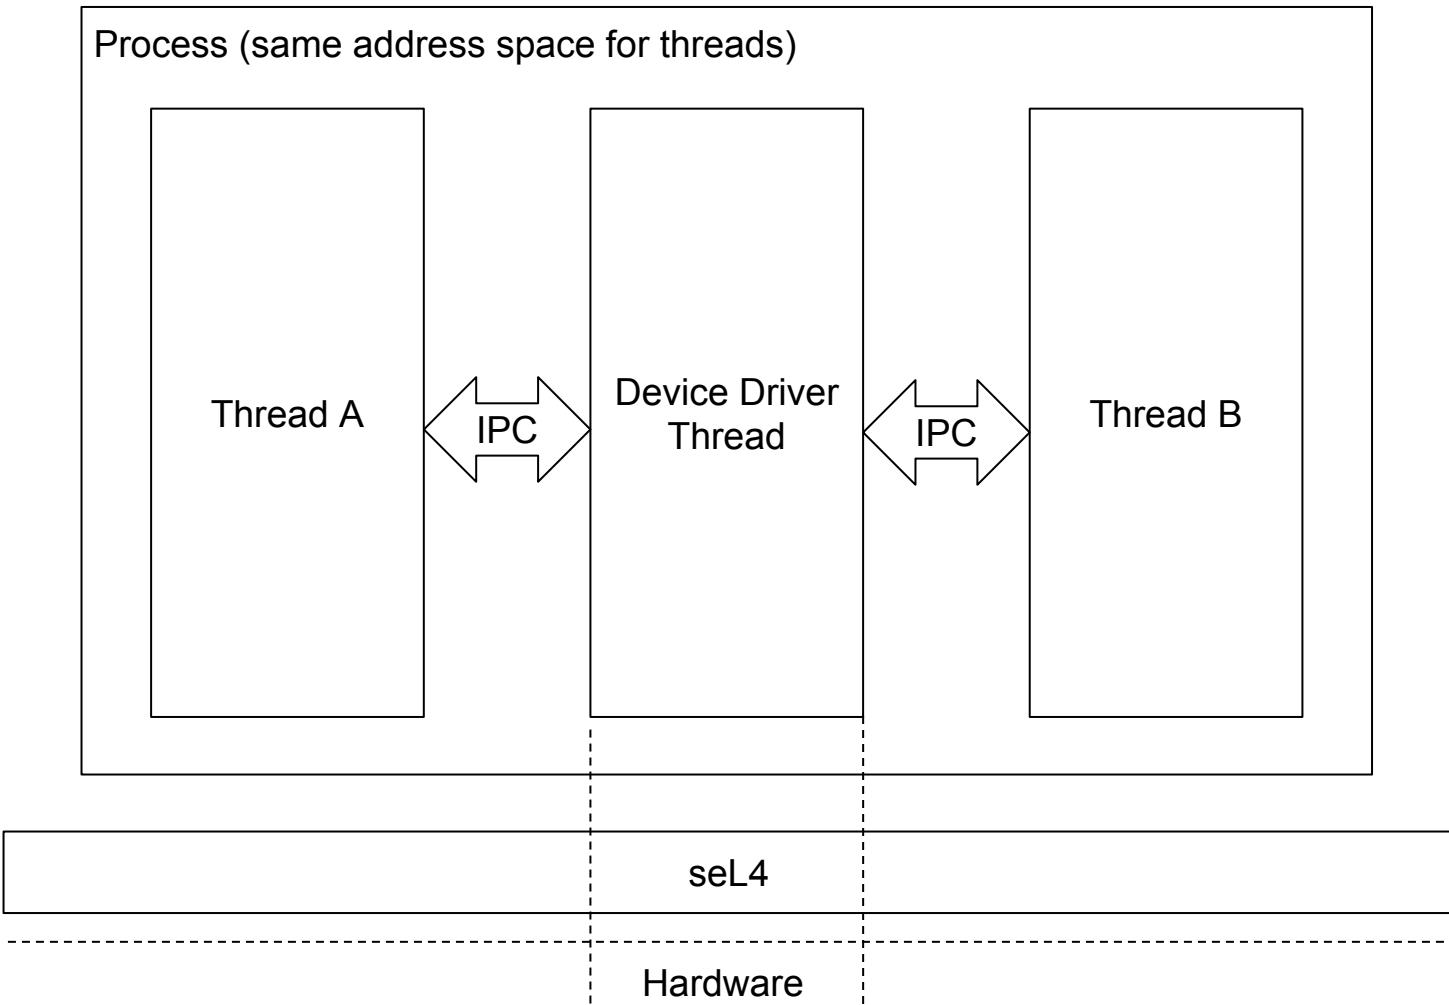
Driver Model



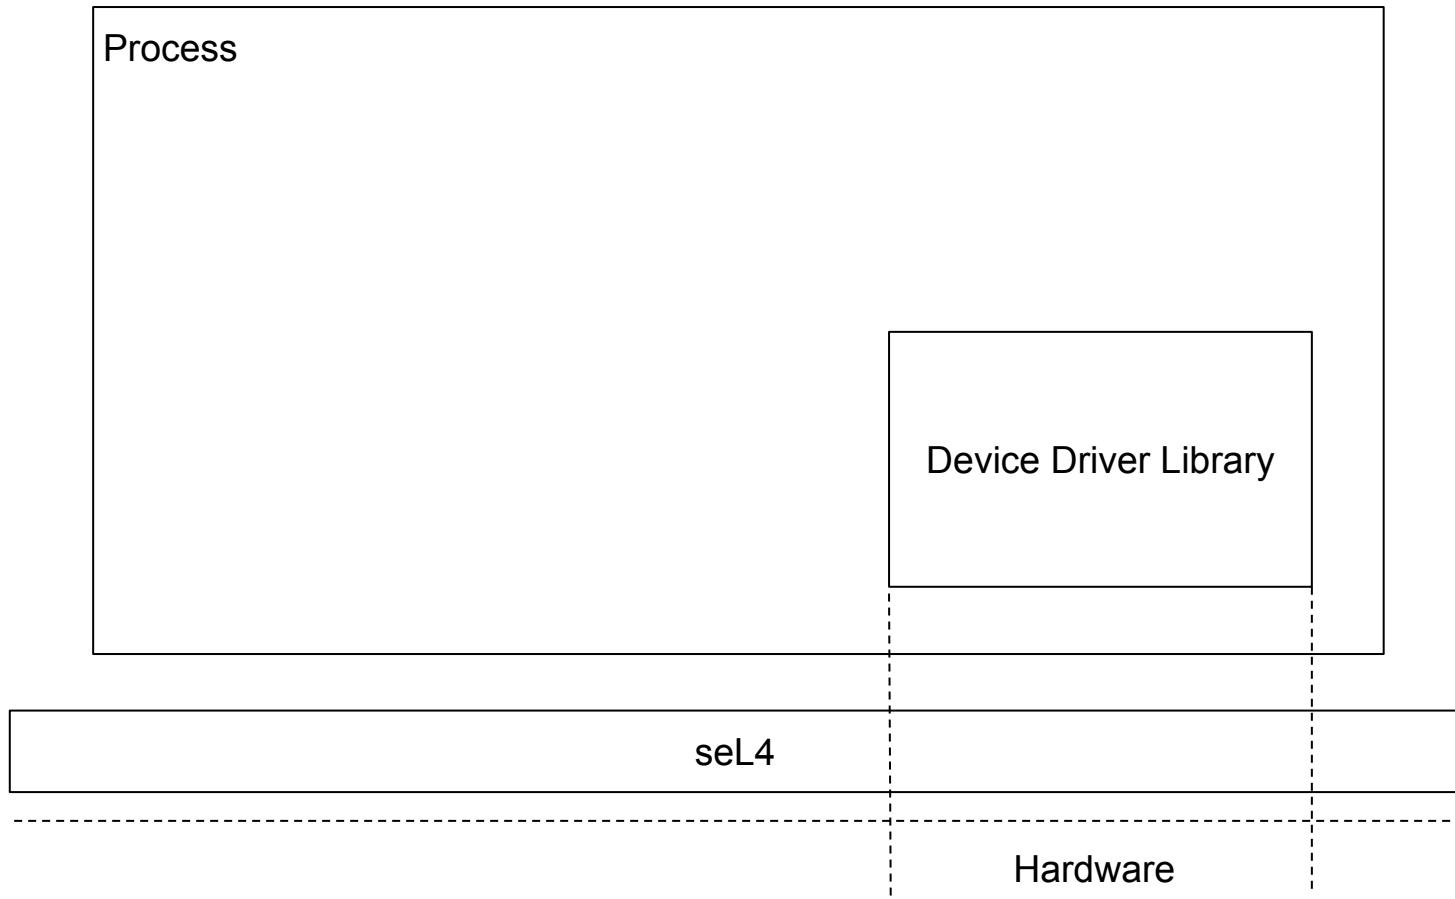
Driver Model



Driver Model



Driver Model



Driver Model



Driver Model

Summary



- Device drivers are at userlevel
- No policy for driver interfaces

Driver Model

Summary



- Device drivers are at userlevel
- No policy for driver interfaces

So how do we access hardware?

Interacting with Hardware



- Memory-Mapped IO
- Interrupts
- IO Ports (x86 only)

Interacting with Hardware

Memory-Mapped IO



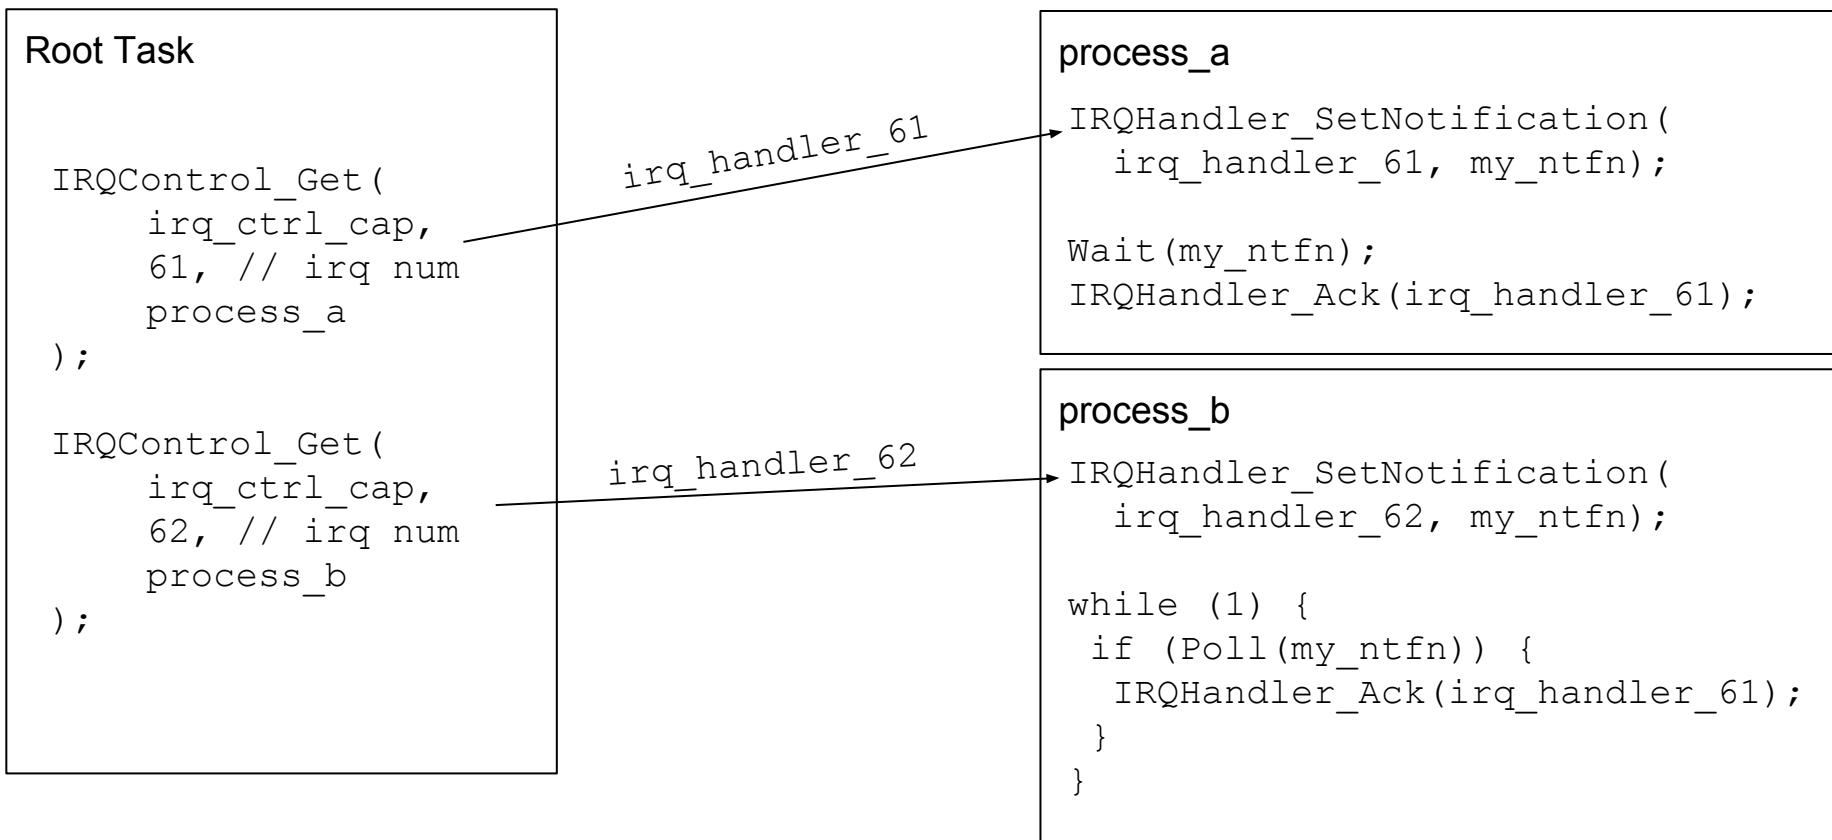
- Kernel creates frames for device regions during boot
- Root task gets caps to all device frames
- `deviceRegions` field in `bootinfo` contains (for each region):
 - `Paddr`
 - `Size`
 - Frame cap `cptrs`
- Map device frames in like any other frame

Interacting with Hardware

Interrupts



- Interrupts will result in signals or notifications
 - So you can `sel4_Wait` or `sel4_Poll` for interrupts



Interacting with Hardware

IO Ports (x86)



- Read and written by invoking io port caps
 - E.g. `sel4_IA32_IOPort_In8(io_port_cap, io_port_num)`
- IO port caps represent access to a contiguous range of IO ports
- IO port caps can be minted into new caps representing a subset of the original cap's range
- Root thread starts with a cap to all IO ports

Interacting with Hardware

Summary



- Device frames are just memory
- Interrupts are delivered as notifications
- Userland has access to IO ports on x86

Up next: Using drivers with seL4

Using Drivers

Libraries



- platsupport
 - seL4-agnostic device drivers
 - Expect memory to be mapped in for them
 - Provide functions for handling interrupts
- sel4platsupport
 - seL4-specific wrappers of platsupport drivers
 - Map in the device frames
 - Receives interrupt notifications and invokes handlers
- CAmkES
 - Generates code to interact with hardware on seL4
 - Abstracts devices as components

Using Drivers

platsupport



- Collections of drivers organised into groups
 - Arch: architecture (ie. arm/x86)
 - Plat: (platform) a specific soc/board (e.g. imx31, exynos4, exynos5)
 - Mach: (machine) a group of platforms with parts in common (e.g. exynos)
- Some devices have common interfaces across platforms
 - E.g. timers, character devices
- Others don't
- Most platforms have drivers for:
 - Serial port
 - Timer
 - Clock
 - i2c
 - gpio

Using Drivers

sel4platsupport



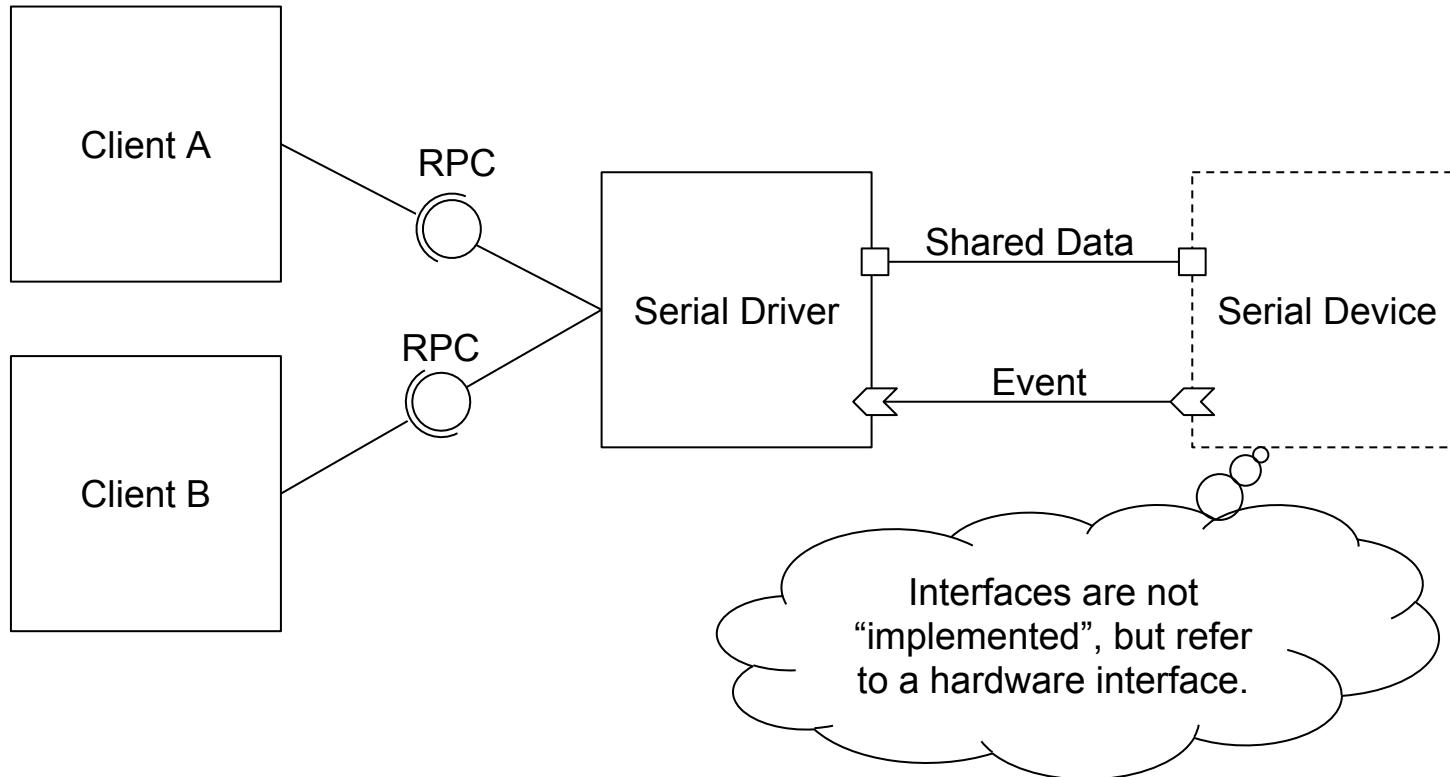
- sel4platsupport drivers need:
 - simple
 - device frame caps and irq control caps
 - vspace
 - mapping device frames
 - vka
 - creating kernel objects
 - managing caps
- Can request driver for default some devices on current platform
 - E.g. `sel4platsupport_get_default_timer`

Using Drivers

CAmkES



- Devices are treated like components



Using Drivers

Accessing Hardware in CAmkES



- Memory-Mapped IO
 - Dataports (shared memory)
- Interrupts
 - Events
- IO Ports (x86)
 - RPC

Drivers talk to devices like any other components!

Using Drivers

Typical Driver Usage



RPC Interface

```
void send(char)    ○—  
char recv()
```

Serial Driver

```
serial_drv_t driver;  
  
void init() {  
    serial_drv_init(&driver, registers);  
}  
  
void driver_send(ch) {  
    driver.send(ch);  
}  
  
char driver_recv() {  
    // generated fn waits for event  
    irq_wait();  
    driver.handle_irq();  
    return driver.recv();  
}
```

Shared Data: “registers”

Event: “irq”

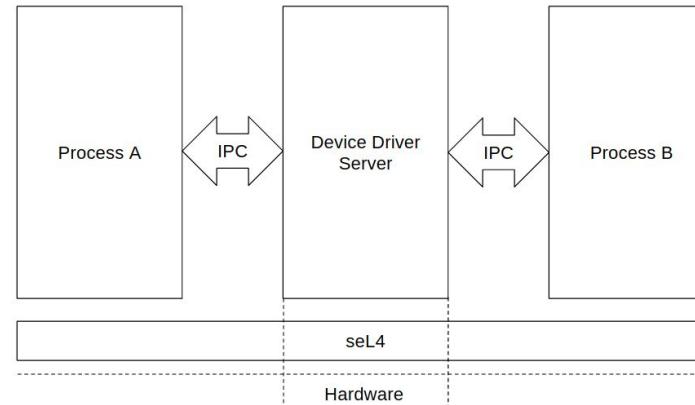
Actual driver implementation lives in platsupport!

Using Drivers

Sound familiar?



Driver Model



3 | Presentation title | Presenter name



Using Drivers

Summary



- Libraries
 - platsupport: seL4 agnostic drivers
 - sel4platsupport: sel4-specific platsupport wrappers
- On bare seL4
 - Provide sel4platsupport drivers with required interfaces
- On CAmkES
 - Devices look like components
 - Driver components are typically wrappers around platsupport drivers

Using Drivers

Exercise



- Initialising/using a platsupport timer driver
 - seL4 tutorials: hello-timer
 - CAmkES tutorials: hello-camkes-timer

Writing Drivers



- A simple driver for an imaginary timer
 - Function to initialise device
 - Functions to perform operations with device]
 - E.g. read/write to a serial port, start/stop/read a timer
 - Functions to handle external events (ie. interrupts)

```
struct tmr_drv {  
    void* vaddr;  
    void (*handle_irq)(tmr_drv* t);  
};  
  
init_timer(tmr_drv* t, void* vaddr)  
{  
    t->vaddr = vaddr;  
    t->handle_irq = my_timer_handle_irq;  
}  
  
my_timer_handle_irq(tmr_drv *t) {  
    // ack the interrupt  
    *(t->vaddr + IRQ_REG_OFFSET) = 0;  
}  
  
start_timer(tmr_drv *t) {  
    *(t->vaddr + CTRL_REG_OFFSET) |=  
        (1 << START_TIMER_BIT);  
}  
  
read_timer_counter(tmr_drv *t) {  
    return *(t->vaddr + COUNT_REG_OFFSET);  
}
```

Writing Drivers



- What resources does this driver need?
 - Someone needs to map the timer device frame into our address space
 - Someone needs to call the “handle_irq” function pointer when a timer interrupt arrives

```
struct tmr_drv {  
    void* vaddr;  
    void (*handle_irq)(tmr_drv* t);  
};  
  
init_timer(tmr_drv* t, void* vaddr)  
{  
    t->vaddr = vaddr;  
    t->handle_irq = my_timer_handle_irq;  
}  
  
my_timer_handle_irq(tmr_drv *t) {  
    // ack the interrupt  
    *(t->vaddr + IRQ_REG_OFFSET) = 0;  
}  
  
start_timer(tmr_drv *t) {  
    *(t->vaddr + CTRL_REG_OFFSET) |=  
        (1 << START_TIMER_BIT);  
}  
  
read_timer_counter(tmr_drv *t) {  
    return *(t->vaddr + COUNT_REG_OFFSET);  
}
```

Writing Drivers

On bare seL4



- Assuming we have:
 - Simple
 - Timer device frame cap
 - Irq handler cap for timer interrupts
 - Vka
 - Vspace

Writing Drivers

On bare seL4



```
// Mapping in the timer device frame

timer_frame_cap = simple_get_frame_cap(&simple, TIMER_PADDR);
timer_vaddr = vspace_map_pages(
    &vspace,
    &timer_frame_cap,      // frame caps
    NULL,                  // cookies
    sel4_CanWrite | sel4_CanRead, // rights
    1,                     // num pages
    12,                    // size bits (4k page)
    0);                   // cacheable

timer_drv t;
init_timer(&t, timer_vaddr);
```

Writing Drivers

On bare seL4



```
// Making interrupt notification

// allocate (retype) notification object
Vka_object_t irq_notification;
vka_alloc_notification(&vka, &irq_notification);

// allocate a cspace slot for irq handler cap
cspacepath_t irq_handler_path;
vka_cspace_alloc_path(&vka, &irq_handler_path);

// get the irq handler cap
simple_get_IRQ_control(&simple, TIMER_IRQ, irq_handler_path);

// associate the irq handler with the notification
sel4_IRQHandler_SetNotification(
    irq_handler_path.capPtr,
    irq_notification.cptr);
```

Writing Drivers

On bare seL4



```
// Invoking callback on interrupt

while (true) {

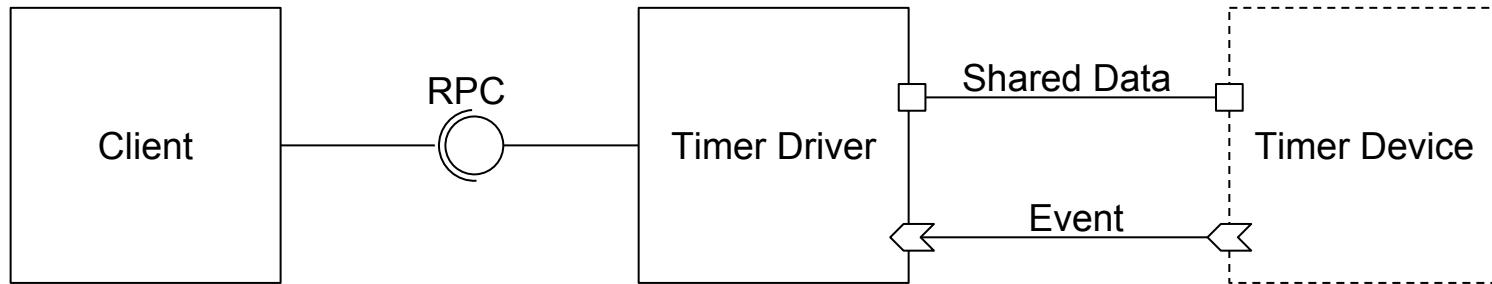
    // block until interrupt arrives
    sel4_WAIT(irq_notification.cptr);

    // invoke callback
    t.handle_irq(&t);

}
```

Writing Drivers

On CAMkES



```
struct tmr_drv {  
    void* vaddr;  
    void (*handle_irq)(tmr_drv* t);  
};  
  
init_timer(tmr_drv* t, void* vaddr)  
{  
    t->vaddr = vaddr;  
    t->handle_irq = my_timer_handle_irq;  
}
```

```
my_timer_handle_irq(tmr_drv *t) {  
    // ack the interrupt  
    *(t->vaddr + IRQ_REG_OFFSET) = 0;  
}  
  
start_timer(tmr_drv *t) {  
    *(t->vaddr + CTRL_REG_OFFSET) |=  
        (1 << START_TIMER_BIT);  
}  
  
read_timer_counter(tmr_drv *t) {  
    return *(t->vaddr + COUNT_REG_OFFSET);  
}
```

Writing Drivers

On CAMkES

```
component Timer {
    hardware;
    dataport Buf regs;
    emits Irq irq;
}

component TimerDriver {
    dataport Buf registers;
    consumes Irq interrupt;
    provides TimerInterface iface;
}

component Client {
    control;
    uses TimerInterface tmr_iface;
}

procedure TimerInterface {
    void start();
    int read();
}
```

```
assembly {
    composition {
        component Timer timer_device;
        component TimerDriver driver;
        component Client client;

        connection seL4HardwareMMIO regs_con(
            from driver.registers,
            to timer_device.regs);
        connection seL4HardwareInterrupt irq_con(
            from timer_device.irq,
            to driver.interrupt);
        connection seL4RPC client_con(
            from client.tmr_iface, to driver.iface);
    } configuration {
        timer_device.regs_attributes =
            "0x12340000:0x1000";
        timer_device.irq_attributes = 42;
    }
}
```



Writing Drivers

On CAMkES

```
// client.c

int run(void) {
    tmr_iface_start();

    while (true) {
        int time = tmr_iface_read();
        printf("%d\n", time);
    }
}
```

Imported from glue code
Exported to glue code

```
// timer_driver.c

tmr_drv timer_driver;

void handle_irq(void) {
    timer_driver.handle_irq(&timer_driver);
    interrupt_reg_callback(handle_irq);
}

void iface_init(void) {
    init_timer(&timer_driver, registers);
    interrupt_reg_callback(handle_irq);
}

void iface_start(void) {
    start_timer(&timer_driver);
}

int iface_read(void) {
    return read_timer_counter(&timer_driver);
}
```



Writing Drivers

Summary



- Drivers broken into 2 parts
 - Interaction with device
 - Agnostic of environment
 - Allocation of system resources
 - Depends on environment
 - CAmkES and sel4 libraries have different approaches

Writing Drivers

Final Exercise



- You should have an emulator for the freescale i.mx31
- Reference manual at <http://goo.gl/MIY9GE>
- Make an app from scratch in CAMkES or bare seL4 (pick your favourite) that prints a message once per second
 - Without using platsupport/sel4platsupport
- Hints
 - Copy and modify a tutorial app
 - Use EPIT2 (frame paddr: 0x53f98000, irq no.: 27)
 - Why not EPIT1?